

SungGeun AN

ssunggun2@snu.ac.kr | ssunggun2.github.io | linkedin.com/in/sunggeun-an | github.com/ssunggun2

RESEARCH INTEREST

- Whether multimodal language models reason from the *evidence in front of them* or from priors learned during pretraining.
- Benchmark and metric design that separates *perception* failures from *reasoning* failures.
- Evaluation methodology for language models, including pass@k diagnostics and human-baseline studies.
- **Keywords:** Multimodal LLMs, physical reasoning, evaluation and benchmark design, retrieval-augmented generation, vision-language models.

EDUCATION

03/2024 – 08/2026	Seoul National University <i>M.S. in Data Science</i> Graduate School of Data Science (GSDS), SKI-ML Lab M.S. thesis, UNREAL: Uncovering Physics Reasoning Limits in MLLMs
2012 – 2021	Seoul National University <i>B.S. in Geography Education</i> College of Education, Department of Geography Education

EXPERIENCE

03/2024 – 08/2026	M.S. Student, Research Assistant <i>Seoul National University, SKI-ML Lab, Seoul, South Korea</i> – M.S. thesis (UNREAL): a benchmark that tests whether multimodal LLMs reason from the video in front of them or fall back on physical priors learned in pretraining. Synthesized 1,528 videos in Kubric under 15 categories of deliberate physics violations with staged QA, and proposed two metrics (Reality Bias Rate, World-Model Inconsistency Rate) validated against a 13-annotator human study – Industry collaboration with LG (2026): led evaluation and dataset research for extending a Korean LLM (EXAONE-3.5-7.8B-Instruct) beyond finance QA into long-tail expert domains, designing a pass@k diagnostic framework (N=128) that separates latent from trained capability and grades benchmarks into difficulty tiers by the pass@1 to pass@128 gap
12/2021 – 12/2023	ML/NLP Engineer <i>Aimpace, Seoul, South Korea</i> – Built and trained Korean language models for address and message extraction – Designed and deployed NLP APIs using FastAPI, serving real-time inference – Implemented OCR pipelines integrating Naver Clova and GCP Vision APIs – Migrated backend from Flask to FastAPI with async processing for 3x throughput – Integrated Redis caching and AWS SQS/Lambda for scalable async workflows – Deployed and managed services on AWS ElasticBeanstalk

PUBLICATIONS

1. **Sieve and Sage: Efficient Distraction Filtering for Reliable RALM Abstention.**
In *Findings of EMNLP*, 2026.
Jongbin Won, **Sung Geun An**, Jay-Yoon Lee.

SERVICE AND ACTIVITIES

02/2026 – 08/2026	Lab manager, SKI-ML Lab, Seoul National University Ran lab operations: intern and member selection and onboarding, seat and resource allocation, and research budget handling with the advisor and administrative staff. Set the usage policy for the shared GPU servers and the LLM API gateway, arbitrating allocation between members and handling outages. Managed the weekly group meeting, paper seminars, and study groups, served as the lab's contact point for external collaborations and department events, and documented the procedures and infrastructure setup so the role could be handed over.
03/2026 – 07/2026	Organizing committee, SNU AI Challenge 2026 Designed the Kaggle task and evaluation metric; scored final reports against a rubric.
03/2025 – 06/2025	Teaching assistant, Machine Learning and Deep Learning 1, SNU GSDS About 100 graduate students. Supported PyTorch implementations of CNN, RNN, and Vision Transformer models, debugged training and environment issues in office hours, and graded assignments and exams against a shared rubric.
09/2024 – 12/2024	Teaching assistant, Foundations of Mathematics and Statistics for Data Science, SNU GSDS About 50 graduate students. Ran review sessions and office hours on linear algebra, multivariate calculus, and probability and statistical inference, and graded assignments and exams against a shared rubric.
08/2025 – 08/2025	Teaching assistant, Samsung Electronics DS2 program NLP and LLM intensive for Samsung Electronics engineers, held at the SNU-Samsung joint research center. Ran TA office-hour sessions spanning the mathematical foundations of deep learning, word embeddings, RNN and LSTM, Transformer encoder-decoder, the Hugging Face stack, prompting and chain-of-thought, LangChain and API usage, retrievers and FAISS-based RAG, and alignment tuning (RLHF, DPO). Resolved environment and pipeline issues during labs and proctored evaluations.
08/2025 – 08/2025	Teaching assistant, Hyundai Motor Company AI training AI and NLP training program for practicing engineers at Hyundai Motor Company. Ran hands-on office-hour sessions, resolved environment and pipeline issues during labs, and proctored evaluations.

SKILLS

Machine Learning & Deep Learning	PyTorch, TensorFlow, Keras, Scikit-learn, Hugging Face Transformers
NLP & Multimodal	Large Language Models, Multimodal LLMs, Vision-Language Models, Retrieval-Augmented Generation, Named Entity Recognition, OCR
Evaluation & Benchmarking	Benchmark Design, Human Evaluation Study Design, pass@k Diagnostics, Prompting Strategy Analysis
Programming	Python, SQL
MLOps & Cloud	AWS (SQS, Lambda, ElasticBeanstalk), Docker, FastAPI, Redis
Languages	Korean (native speaker), English (professional working proficiency)